

# 上海复旦大学物理楼 324 机房

## 测试集群概况介绍

### 一、集群信息介绍

1. 此次测试集群包含 1 台登陆节点、1 台存储节点也是管理节点、10 台刀片计算节点，共计 12 台主机组成。
2. 网络系统包括 ssh 千兆管理网和 InfiniBand 40Gb 存储网。
3. 管理节点作为存储节点、PBS 调度服务端、NIS 用户认证服务端，是真正意义上的管理节点。
4. 集群采用 PBS 作业调试系统。
5. 主机名查看：

```
[root@login ~]# cat /etc/hosts
127.0.0.1 localhost localhost.localdomain localhost4 localhost4.localdomain4
::1 localhost localhost.localdomain localhost6 localhost6.localdomain6
#####--SSH-MGMT--#####
11.11.11.254 admin node254
11.11.11.253 login node253
11.11.11.1 node1
11.11.11.2 node2
11.11.11.3 node3
11.11.11.4 node4
11.11.11.5 node5
11.11.11.6 node6
11.11.11.7 node7
11.11.11.8 node8
11.11.11.9 node9
11.11.11.10 node10
#####--IB-STORAGE--#####
12.12.12.254 iadmin
12.12.12.253 ilogin
12.12.12.1 inode1
12.12.12.2 inode2
12.12.12.3 inode3
12.12.12.4 inode4
12.12.12.5 inode5
12.12.12.6 inode6
12.12.12.7 inode7
12.12.12.8 inode8
12.12.12.9 inode9
12.12.12.10 inode10
```

## 二、登陆集群

login 已接入校园网，IP 地址为 10.92.3.111。

```
[root@login ~]# ip a
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
    inet6 ::1/128 scope host
        valid_lft forever preferred_lft forever
2: eth0: <BROADCAST,MULTICAST> mtu 1500 qdisc pfifo_fast state DOWN qlen 1000
    link/ether 00:1b:21:89:0a:a6 brd ff:ff:ff:ff:ff:ff
3: eth1: <BROADCAST,MULTICAST> mtu 1500 qdisc pfifo_fast state DOWN qlen 1000
    link/ether 00:1b:21:89:0a:a7 brd ff:ff:ff:ff:ff:ff
4: eth2: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc pfifo_fast state UP qlen 1000
    link/ether 00:25:90:34:67:da brd ff:ff:ff:ff:ff:ff
    inet 11.11.11.253/16 brd 11.11.255.255 scope global eth2
    inet6 fe80::225:90ff:fe34:67da/64 scope link
        valid_lft forever preferred_lft forever
5: eth3: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc pfifo_fast state UP qlen 1000
    link/ether 00:25:90:34:67:db brd ff:ff:ff:ff:ff:ff
    inet 10.92.3.111/24 brd 10.92.3.255 scope global eth3
    inet6 2001:da8:8001:9203:225:90ff:fe34:67db/64 scope global dynamic
        valid_lft 2591869sec preferred_lft 2591869sec
    inet6 fe80::225:90ff:fe34:67db/64 scope link
        valid_lft forever preferred_lft forever
6: ib0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 2044 qdisc mq state UP qlen 1024
    link/infiniband a0:00:02:20:fe:80:00:00:00:00:00:00:00:00:02:c9:03:00:0f:61:0b brd 00:f
    inet 12.12.12.253/24 brd 12.12.12.255 scope global ib0
    inet6 fe80::202:c903:f:610b/64 scope link
        valid_lft forever preferred_lft forever
```

## 三、账户管理

1. 集群采用 NIS 进行统一用户认证，创建用户需要以管理员身份 ssh 到 admin 节点操作，示例创建 test 用户；  
`useradd test -d /public/home/test`
2. 为新用户创建密码  
`passwd test`
3. 同步用户账户到所有节点（此 3 步操作已完成新增用户操作）；  
`make -C /var/yp`
4. 用户可自行修改密码，直接在 login 上操作即可；  
`yppasswd`

## 四、普通用户测试 PBS 作业系统可用性

1. 调度系统概况查询；

```
[ll@admin ~]$ pestat
node      state  load    phymem  ncpus   allmem  resi   usrs   tasks  jobidlist
node10    free   2.24    24020   12      56788  790    2/1    2      8 ll
node1     free   2.33    24020   12      56788  811    2/1    2      8 ll
node2     free   2.22    24020   12      56788  787    2/1    2      8 ll
node3     free   2.26    19980   12      52748  741    2/1    2      8 ll
node4     free   2.03    24020   12      56788  781    2/1    2      8 ll
node5     free   2.31    24020   12      56788  786    2/1    2      8 ll
node6     free   2.80*   24020   12      56788  820    3/1    2      8 ll
node7     free   2.27    24020   12      56788  791    2/1    2      8 ll
node8     free   2.33    24020   12      56788  794    2/1    2      8 ll
node9     free   2.19    24020   12      56788  788    2/1    2      8 ll
```

2. 编写作业脚本;

```
[ll@login ~]$ cat iozone.pbs
#!/bin/bash
#PBS -N iozone_storage
#PBS -l nodes=10:ppn=2
#PBS -j oe
#PBS -l walltime=24:00:00
#PBS -q middle

#####
cd /public/software/benchmark/iozone/gnu/3.430/
RSH=ssh ./iozone -i 0 -i 1 -r 1M -s 48g -t 20 -+m /public/home/ll/test/nodelist
```

3. 提交作业 sbatch iozone.pbs。

4. 查询作业正常运行;

```
[ll@admin ~]$ qstat
Job ID          Name          User          Time Use S Queue
-----
8.admin         iozone_storage ll             0 R middle
[ll@admin ~]$ qstat -n 8

admin:

Job ID          Username      Queue          Jobname          SessID  NDS   TSK   Req'd  Req'd   Elap
Memory         Time        S          Time
-----
8.admin         ll           middle         iozone_storage   0       10    20    --    24:00:00 R  00:00:18
node6/0+node6/1+node9/0+node9/1+node4/0+node4/1+node5/0+node5/1+node2/0
+node2/1+node7/0+node7/1+node8/0+node8/1+node10/0+node10/1+node1/0+node1/1
+node3/0+node3/1
[ll@admin ~]$
```

5. 确认计算节点有 iozone 测试进程在运行;

```
[node8]ll      923    1 0 20:06 ?      00:00:11 /public/software/benchmark/iozone/gnu/3.430/iozone --s -t 1 -r 4 -s 4 --c node6
[node8]ll      924    923 0 20:06 ?      00:00:00 /public/software/benchmark/iozone/gnu/3.430/iozone --s -t 1 -r 4 -s 4 --c node6
[node8]root    939    2 0 Jan05 ?      00:00:00 [edac-poller]
[node8]root    3550   3548 0 20:48 ?      00:00:00 bash -c PATH=/sbin:/usr/sbin:/bin:/usr/bin ps -ef | grep ll
[node8]root    3593   3550 0 20:48 ?      00:00:00 grep ll
[node8]root    16370  1 0 Jan06 ?      00:01:04 /opt/gridview/gvsnmp/sbin/snmpd -c /opt/gridview/gvsnmp/conf/snmpd.conf -M /opt/g
[node10]root    595    2 0 Jan05 ?      00:00:03 [mpt_poll_0]
[node10]ll      935    1 0 20:06 ?      00:00:15 /public/software/benchmark/iozone/gnu/3.430/iozone --s -t 1 -r 4 -s 4 --c node6
[node10]root    936    2 0 Jan05 ?      00:00:00 [edac-poller]
[node10]ll      937    935 0 20:06 ?      00:00:00 /public/software/benchmark/iozone/gnu/3.430/iozone --s -t 1 -r 4 -s 4 --c node6
[node10]ll      964    1 0 20:06 ?      00:00:15 /public/software/benchmark/iozone/gnu/3.430/iozone --s -t 1 -r 4 -s 4 --c node6
[node10]ll      965    964 0 20:06 ?      00:00:00 /public/software/benchmark/iozone/gnu/3.430/iozone --s -t 1 -r 4 -s 4 --c node6
[node10]root    4028   4025 0 20:48 ?      00:00:00 bash -c PATH=/sbin:/usr/sbin:/bin:/usr/bin ps -ef | grep ll
[node10]root    4070   4028 0 20:48 ?      00:00:00 grep ll
[node10]root    16391  1 0 Jan06 ?      00:01:04 /opt/gridview/gvsnmp/sbin/snmpd -c /opt/gridview/gvsnmp/conf/snmpd.conf -M /opt/g
d
[node3]root     581    2 0 Jan05 ?      00:00:03 [mpt_poll_0]
[node3]root     869    2 0 Jan05 ?      00:00:00 [edac-poller]
[node3]ll      5398    1 0 20:06 ?      00:00:15 /public/software/benchmark/iozone/gnu/3.430/iozone --s -t 1 -r 4 -s 4 --c node6
[node3]ll      5399   5398 0 20:06 ?      00:00:00 /public/software/benchmark/iozone/gnu/3.430/iozone --s -t 1 -r 4 -s 4 --c node6
[node3]ll      5425    1 0 20:06 ?      00:00:15 /public/software/benchmark/iozone/gnu/3.430/iozone --s -t 1 -r 4 -s 4 --c node6
[node3]ll      5426   5425 0 20:06 ?      00:00:00 /public/software/benchmark/iozone/gnu/3.430/iozone --s -t 1 -r 4 -s 4 --c node6
[node3]root    8461   8459 0 20:48 ?      00:00:00 bash -c PATH=/sbin:/usr/sbin:/bin:/usr/bin ps -ef | grep ll
[node3]root    8499   8461 0 20:48 ?      00:00:00 grep ll
[node3]root    11643  1 0 Jan06 ?      00:02:46 /opt/gridview/gvsnmp/sbin/snmpd -c /opt/gridview/gvsnmp/conf/snmpd.conf -M /opt/g
[node2]root     503    2 0 Jan05 ?      00:00:03 [mpt_poll_0]
[node2]root    1052    2 0 Jan05 ?      00:00:00 [edac-poller]
[node2]ll      5779    1 0 20:06 ?      00:00:14 /public/software/benchmark/iozone/gnu/3.430/iozone --s -t 1 -r 4 -s 4 --c node6
[node2]ll      5780   5779 0 20:06 ?      00:00:00 /public/software/benchmark/iozone/gnu/3.430/iozone --s -t 1 -r 4 -s 4 --c node6
[node2]ll      5806    1 0 20:06 ?      00:00:14 /public/software/benchmark/iozone/gnu/3.430/iozone --s -t 1 -r 4 -s 4 --c node6
[node2]ll      5807   5806 0 20:06 ?      00:00:00 /public/software/benchmark/iozone/gnu/3.430/iozone --s -t 1 -r 4 -s 4 --c node6
[node2]root    8771   8769 0 20:48 ?      00:00:00 bash -c PATH=/sbin:/usr/sbin:/bin:/usr/bin ps -ef | grep ll
[node2]root    8814   8771 0 20:48 ?      00:00:00 grep ll
[node2]root    13317  1 0 Jan06 ?      00:02:39 /opt/gridview/gvsnmp/sbin/snmpd -c /opt/gridview/gvsnmp/conf/snmpd.conf -M /opt/g
[node5]root     580    2 0 Jan05 ?      00:00:03 [mpt_poll_0]
[node5]root     942    2 0 Jan05 ?      00:00:00 [edac-poller]
[node5]ll      5513    1 0 20:06 ?      00:00:15 /public/software/benchmark/iozone/gnu/3.430/iozone --s -t 1 -r 4 -s 4 --c node6
[node5]ll      5514   5513 0 20:06 ?      00:00:00 /public/software/benchmark/iozone/gnu/3.430/iozone --s -t 1 -r 4 -s 4 --c node6
[node5]ll      5540    1 0 20:06 ?      00:00:15 /public/software/benchmark/iozone/gnu/3.430/iozone --s -t 1 -r 4 -s 4 --c node6
[node5]ll      5541   5540 0 20:06 ?      00:00:00 /public/software/benchmark/iozone/gnu/3.430/iozone --s -t 1 -r 4 -s 4 --c node6
[node5]root    8515   8513 0 20:48 ?      00:00:00 bash -c PATH=/sbin:/usr/sbin:/bin:/usr/bin ps -ef | grep ll
[node5]root    8561   8515 0 20:48 ?      00:00:00 grep ll
[node5]root    13444  1 0 Jan06 ?      00:02:49 /opt/gridview/gvsnmp/sbin/snmpd -c /opt/gridview/gvsnmp/conf/snmpd.conf -M /opt/g
[node7]root     494    2 0 20:48 ?      00:00:00 bash -c PATH=/sbin:/usr/sbin:/bin:/usr/bin ps -ef | grep ll
[node7]root     547    494 0 20:48 ?      00:00:00 grep ll
[node7]root     580    2 0 Jan05 ?      00:00:03 [mpt_poll_0]
[node7]root     938    2 0 Jan05 ?      00:00:00 [edac-poller]
```

(结尾)